

# *High Performance Computing*

SAGAR BHATT

Person Number: 50170651

Department of Mechanical and Aerospace Engineering,  
University at Buffalo

## I. MATRIX ADDITION PERFORMANCE STUDY

I. Profiling performance using gprof:

**Looping over rows first:**

Flat profile:

Each sample counts as 0.01 seconds.

% time	cumulative seconds	self seconds	calls	self s/call	total s/call	name
61.32	2.44	2.44	1	2.44	2.44	main
38.70	3.98	1.54				__intel_ssse3_rep_memmove
0.00	3.98	0.00	30000	0.00	0.00	std::vector<double, std::allocator<double> >::vector(std::vector<double, std::allocator<double> > const&)
0.00	3.98	0.00	3	0.00	0.00	std::vector<std::vector<double, std::allocator<double> > >::vector(unsigned long, std::allocator<std::vector<double, std::allocator<double> > > const&)
0.00	3.98	0.00	3	0.00	0.00	std::vector<double, std::allocator<double> >::vector(unsigned long, std::allocator<double> > const&)
0.00	3.98	0.00	2	0.00	0.00	get_time
0.00	3.98	0.00	1	0.00	0.00	diff

**Looping over columns first:**

Flat profile:

Each sample counts as 0.01 seconds.

% time	cumulative seconds	self seconds	calls	self s/call	total s/call	name
91.68	16.58	16.58	1	16.58	16.58	main
8.35	18.09	1.51				__intel_ssse3_rep_memmove
0.00	18.09	0.00	30000	0.00	0.00	std::vector<double, std::allocator<double> >::vector(std::vector<double, std::allocator<double> > const&)
0.00	18.09	0.00	3	0.00	0.00	std::vector<std::vector<double, std::allocator<double> > >::vector(unsigned long, std::allocator<std::vector<double, std::allocator<double> > > const&)
0.00	18.09	0.00	3	0.00	0.00	std::vector<double, std::allocator<double> >::vector(unsigned long, std::allocator<double> > const&)
0.00	18.09	0.00	2	0.00	0.00	get_time
0.00	18.09	0.00	1	0.00	0.00	diff

## II. Instrumenting the code with PAPI

### Looping over rows first:

Time taken when looping over rows first:0.850615

The total instructions executed for addition are 600666184

The total cycles used are 1247329260

The total cache misses are 37982311

Real\_time: 5.974195

Proc\_time: 5.960402

Total flops: 303715595

MFLOPS: 50.955555

### Looping over columns first:

Time taken when looping over columns first:14.9937

The total instructions executed for addition are 1000861067

The total cycles used are 30122730315

The total cache misses are 723035278

Real\_time: 18.859709

Proc\_time: 18.812635

Total flops: 300962949

MFLOPS: 15.997914

## II. MANDLEBROT SET

Estimated area of the Mandlebrot set: 1.5095

Resolution:  $\frac{4}{1000}$  for comparing different schedulings. For the plot, resolution was  $\frac{4}{10000}$

Iterations: 10000

*Real Axis<sub>max</sub>* = 2

*Real Axis<sub>min</sub>* = -2

*Imaginary Axis<sub>max</sub>* = 2

*Imaginary Axis<sub>min</sub>* = -2

### Code Performance:


Time without omp: 5.99114 s

Area of the mandlebrot set =1.5095


Read successfully

The total instructions executed are 20876783386  
The total cycles used are 14252196016  
The total cache misses are 274536

**Plot:**



**Figure 1:** *Plot of Mandelbrot Set*



**Figure 2:** *Fringes of the mandlebrot set*

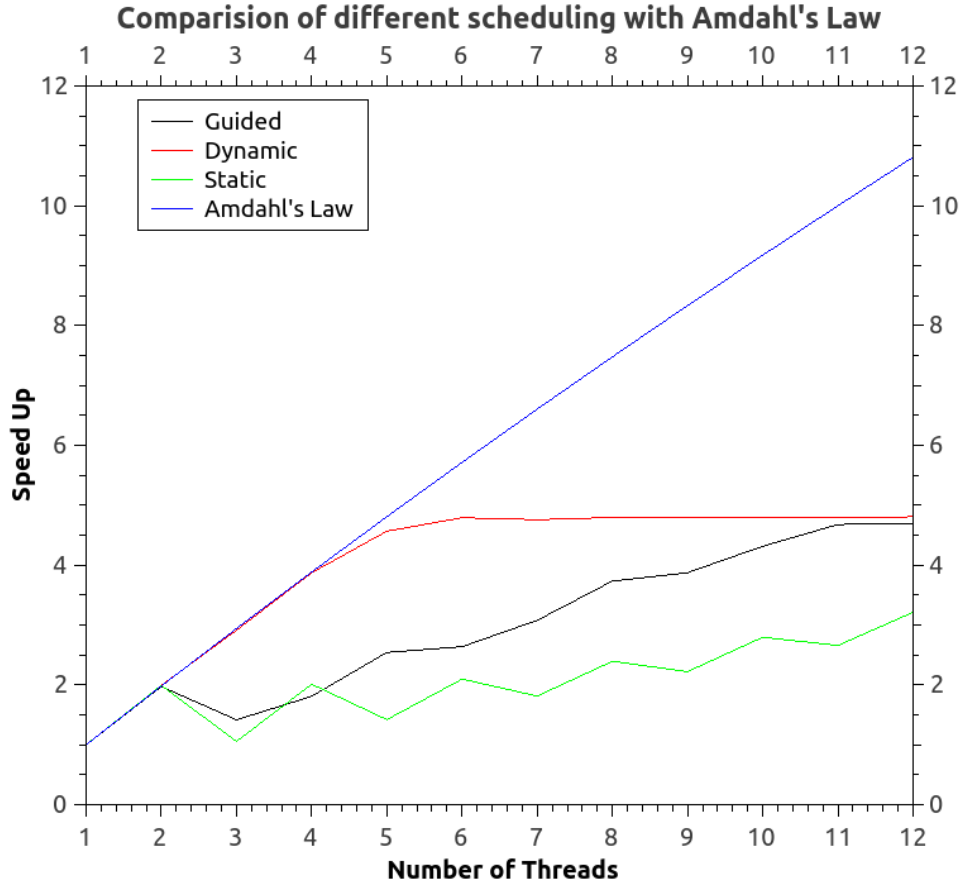
Speedup was computed using the formula:

$$s = \frac{T(1)}{T(p)}$$

The serial fractions were computed using the Karp-Flatt metric:  $f = \frac{1/s - 1/p}{1 - 1/p}$

**Serial fractions for different scheduling:**

Threads	Static	Dynamic	Guided
1	NA	NA	NA
2	0.502304	0.502482	0.507148
3	0.947342	0.34342	0.706208
4	0.497382	0.25828	0.552642
5	0.70475	0.219027	0.393479
6	0.477003	0.208744	0.379274
7	0.551515	0.210188	0.324986
8	0.417825	0.208535	0.267883
9	0.450383	0.208958	0.258204
10	0.357674	0.208973	0.231806
11	0.376339	0.208549	0.213792
12	0.311293	0.208032	0.212929



**Figure 3:** Comparison of scheduling with Amdahl's law

### III. GOLDBACH CONJECTURE

A code was written to verify that Goldbach conjecture was true upto any arbitrary number of integers. For this study, the code was used to verify the conjecture for all even numbers less than 1000.

#### Code Performance:

The total instructions executed are 51546  
The total cycles used are 126342  
The total cache misses are 1277  
Total flops: 672  
MFLOPS: 2.507463

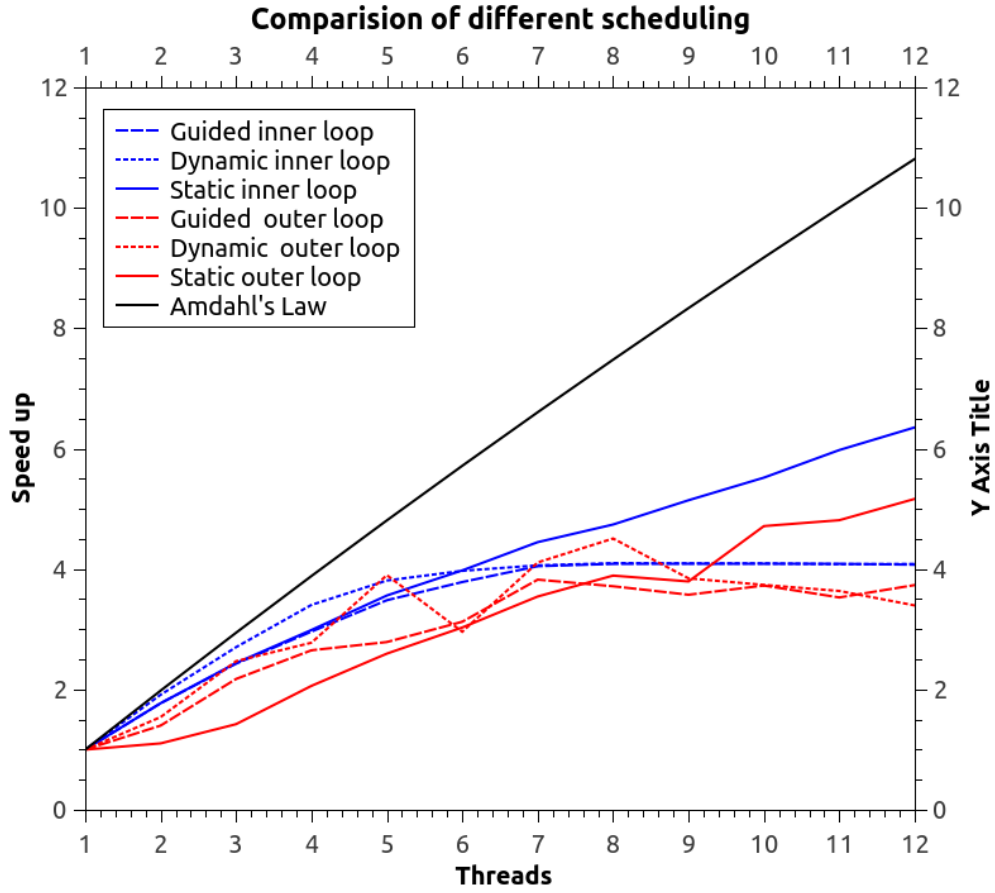


Figure 4: Comparison of scheduling with Amdahl's law

Serial fractions for different scheduling:

PARALLELIZING OUTER LOOP:				PARALLELIZING INNER LOOP:			
Threads	Static	Dynamic	Guided	Threads	Static	Dynamic	Guided
1	NA	NA	NA	1	NA	NA	NA
2	0.906107	0.648986	0.714653	2	0.564533	0.523178	0.564874
3	0.702961	0.405036	0.460048	3	0.410179	0.370116	0.411896
4	0.485634	0.360256	0.377353	4	0.334757	0.293796	0.33783
5	0.385431	0.256561	0.358821	5	0.280915	0.262536	0.287219
6	0.330181	0.338079	0.320094	6	0.251457	0.251904	0.264308
7	0.282251	0.243374	0.261548	7	0.224929	0.246262	0.247159
8	0.257039	0.221895	0.269296	8	0.211026	0.244044	0.244788
9	0.263636	0.260092	0.279905	9	0.194478	0.244254	0.244667
10	0.212191	0.267685	0.268455	10	0.181242	0.244281	0.244685
11	0.207875	0.275105	0.283498	11	0.167339	0.245056	0.244719
12	0.193629	0.294559	0.26786	12	0.157428	0.244818	0.245595

INFERENCE: It is evident from the aforementioned results, that parallelizing inner loop has resulted in a much lower serial fraction hence a much better speed up. This can be also observed in the plot where the general trend for all schedulings shows that speed up is higher when inner loop is parallelized.

#### IV. CALCULATING $\pi$

A program was written to compute  $\pi$  using midpoint rule.

##### **Result:**

Static Scheduling:

Pi=3.14159305347414097

Time taken with static scheduling on 12 threads: 0.510051985000000041 s

Dynamic Scheduling:

Pi=3.14159305347418005

Time taken with dynamic scheduling on 12 threads: 0.5880544349999999986 s

Guided Scheduling:

Pi=3.14159305347416096

Time taken with guided scheduling on 12 threads: 0.5650294710000000005 s

##### **Code Performance:**

The total instructions executed are 12000071835

The total cycles used are 54511050117

The total cache misses are 4035

Total flops: 9500601275

MFLOPS: 415.487762

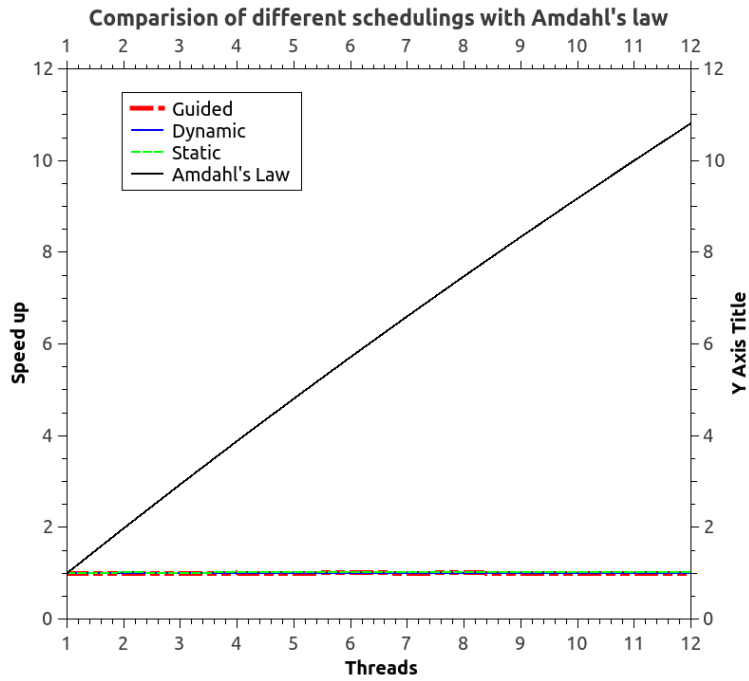


Figure 5: Comparison of scheduling with Amdahl's law

The above plot does not show a lot of difference between the schedulings since the speedup was not good enough. The following plot shows a comparison among different strategies :

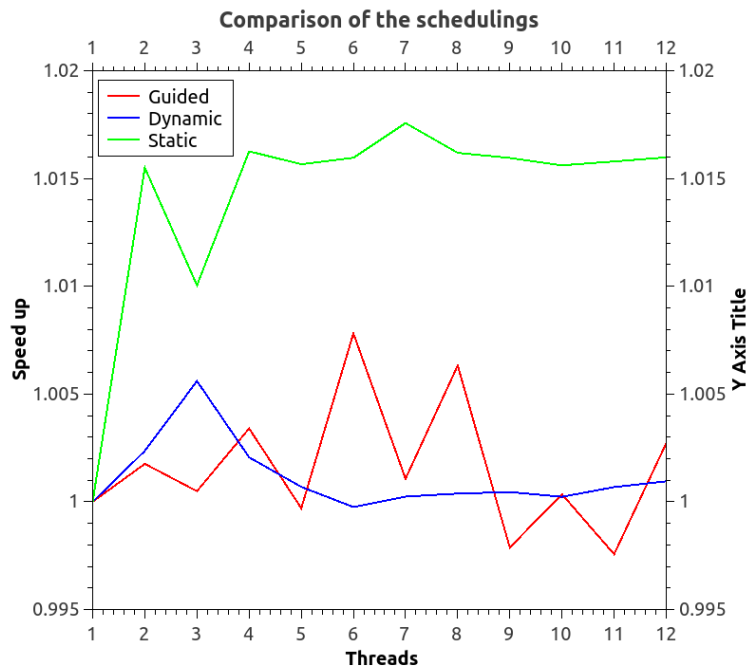


Figure 6: Comparison of scheduling



Although there was not much of a difference between different schedulings, the above plot shows that there was a little (though insignificant) better speedup with static scheduling.

**Serial fractions for different scheduling:**

Threads	Static	Dynamic	Guided
1	NA	NA	NA
2	0.984737	0.997641	0.998241
3	0.99006	0.994427	0.999511
4	0.983998	0.99794	0.996614
5	0.984577	0.999306	1.0003
6	0.984287	1.00024	0.992252
7	0.982735	0.999757	0.998932
8	0.984064	0.99962	0.993718
9	0.984305	0.999552	1.00212
10	0.984625	0.999767	0.999647
11	0.984454	0.999316	1.00244
12	0.984267	0.999049	0.997283